

Exercises week 5 report

Exercise 5.1

1. Use Benchmark.java to run the Mark1 through Mark6 measurements.

- Mark1

```
zq2lii@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:29:09+0200
0.068 s    3.4ns

BUILD SUCCESSFUL in 1s
3 actionable tasks: 2 executed, 1 up-to-date
```

- Mark2

```
zq2lii@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:29:57+0200
24.5 ns

BUILD SUCCESSFUL in 4s
3 actionable tasks: 2 executed, 1 up-to-date
```

- Mark3

```
zq2lii@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:30:31+0200
24.5 ns
26.4 ns
29.1 ns
25.9 ns
25.6 ns
30.4 ns
25.7 ns
25.4 ns
25.5 ns
25.4 ns

BUILD SUCCESSFUL in 27s
3 actionable tasks: 2 executed, 1 up-to-date
```

- Mark4

```

● zq2lii@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:31:30+0200
        25.6 ns +/-  0.260

BUILD SUCCESSFUL in 27s
3 actionable tasks: 2 executed, 1 up-to-date

```

◦ Mark5

```

● zq2lii@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:32:35+0200
488.5 ns +/-  722.55      2
212.9 ns +/-   75.21      4
179.3 ns +/-   49.16      8
167.2 ns +/-   84.11     16
 48.7 ns +/-    5.83     32
 46.7 ns +/-    8.12     64
 58.8 ns +/-   90.06    128
 32.7 ns +/-    7.86    256
 29.9 ns +/-    5.31    512
114.8 ns +/-  200.07   1024
 27.3 ns +/-    0.37   2048
 30.7 ns +/-    8.12   4096
 25.8 ns +/-    1.01   8192
 61.6 ns +/-  104.96  16384
 30.6 ns +/-    7.99  32768
 31.2 ns +/-   10.26  65536
 28.3 ns +/-    3.09  131072
 26.1 ns +/-    0.99  262144
 25.5 ns +/-    0.18  524288
 25.6 ns +/-    0.24 1048576
 25.5 ns +/-    0.21 2097152
 25.7 ns +/-    0.35 4194304
 25.7 ns +/-    0.19 8388608
 26.0 ns +/-    0.75 16777216

BUILD SUCCESSFUL in 10s
3 actionable tasks: 2 executed, 1 up-to-date

```

◦ Mark6

```
zq2lll@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:33:23+0200
multiply      579.2 ns      941.92      2
multiply      260.7 ns      78.61      4
multiply      243.7 ns      60.83      8
multiply      329.2 ns      47.12     16
multiply      346.1 ns      96.26     32
multiply      314.7 ns      20.90     64
multiply      130.7 ns     243.94    128
multiply       59.4 ns       3.77    256
multiply       49.1 ns      18.60    512
multiply       60.5 ns      11.48   1024
multiply       73.8 ns      35.12   2048
multiply       50.2 ns      21.21   4096
multiply       32.5 ns       2.89   8192
multiply       40.3 ns      21.26  16384
multiply       36.1 ns      21.45  32768
multiply       29.4 ns       4.14  65536
multiply       39.0 ns      13.22 131072
multiply       35.3 ns      11.68 262144
multiply       35.1 ns       7.08 524288
multiply       30.9 ns       9.02 1048576
multiply       26.6 ns       1.21 2097152
multiply       26.7 ns       1.79 4194304
multiply       27.7 ns       2.39 8388608
multiply       29.9 ns       7.07 16777216

BUILD SUCCESSFUL in 10s
3 actionable tasks: 2 executed, 1 up-to-date
```

All the results above are plausible. The actual running time may vary from the *Microbenchmarks note*, however they're reasonable.

The running time in Mark1 is significantly small, this may be because in JIT, the loop isn't processed actually. Once we return *dummy*, the result in Mark2 is more reasonable. And in Mark3 we can see all the iterations take nearly the same amount of time.

There're some exception in Mark5 when there are sudden increase for iteration count 128/1024/16384, but we also see that the standard deviations are very large in those cases, which tells us we can have no confidence in those results.

2. Use Mark7 to measure the execution time for the mathematical functions `pow`, `exp`, and so on, as in *Microbenchmarks note* Section 4.2. Record the results in a text file along with appropriate system identification. Preferably do this on at least two different platforms, eg. your own computer and a fellow student/friends computer.

- MacOS

```

• zq2l1i@hostmac code-exercises % gradle -PmainClass=exercises05.Benchmark run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T22:56:04+0200
pow          21.0 ns      2.73   16777216
exp          21.7 ns      0.05   16777216
log          21.6 ns      0.05   16777216
sin          24.9 ns      0.10   16777216
cos          25.2 ns      0.05   16777216
tan          28.9 ns      0.09   16777216
asin        98.5 ns      1.58   4194304
acos       103.3 ns      1.46   4194304
atan        26.8 ns      0.07   16777216

BUILD SUCCESSFUL in 1m 15s
3 actionable tasks: 2 executed, 1 up-to-date

```

- Windows

some analysis here...

Exercise 5.2

1. First compile and run the thread timing code as is, using Mark6, to get a feeling for the variation and robustness of the results. Do not hand in the results but discuss any strangenesses, such as large variation in the time measurements for each cases.

There're always some exceptions in different cases. Such as for *Thread create start join*, when iteration count is 512, there's significant increase of the mean time. This outlier measurement may be caused by the garbage collector accidentally performing some work at that time, or the just-in-time compiler, or some other external disturbance.

2. Now change all the measurements to use Mark7, which reports only the final result. Record the results in a text file along with appropriate system identification.

```

• zq2l1i@hostmac code-exercises % gradle -PmainClass=exercises05.TestTimeThreads run

> Task :app:run
# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Homebrew; 18.0.2.1
# CPU:   null; 4 "cores"
# Date:  2022-10-13T23:11:33+0200
Mark 6 measurements
hashCode()          2.7 ns      0.06   134217728
Point creation      43.3 ns      1.26   8388608
Thread's work       5423.0 ns    120.91   65536
Thread create       856.5 ns      5.47   524288
Thread create start 69050.4 ns   6118.41   4096
Thread create start join 81931.2 ns   6067.93   4096
ai value = 1474500000
Uncontended lock    18.8 ns      0.07   16777216

BUILD SUCCESSFUL in 50s
3 actionable tasks: 2 executed, 1 up-to-date

```

Result is plausible, we can see that the creation of simple object cost just 2.7ns, but the creation of thread takes more than 800ns. And the start of a thread is even more, it takes almost 70000ns, even after creating those threads.

Exercise 5.3

1. Measure the performance of the primecounting example on your own hardware, as a function of the number of threads used to determine whether a given number is a prime. Record system information as well as the measurement results for 1. . . 32 threads in a text file. If the measurements take excessively long time on your computer, you may measure just for 1. . . 16 threads instead.
2. Reflect and comment on the results; are they plausible? Is there any reasonable relation between the number of threads that gave best performance, and the number of cores in the computer you ran the benchmarks on? Any surprises

```
# OS: Mac OS X; 12.3.1; x86_64
# JVM: Oracle Corporation; 17.0.3
# CPU: null; 4 "cores"
# Date: 2022-10-13T23:27:50+0200
countSequential          9807341.5 ns  243944.92      32
countParallelN           1      8542845.8 ns  24560.80      32
countParallelN           2      5439761.8 ns  49421.52      64
countParallelN           3      6450739.3 ns  333770.48     64
countParallelN           4      5796062.9 ns  46225.90      64
countParallelN           5      6153400.0 ns  109422.57     64
countParallelN           6      5728994.7 ns  56877.33      64
countParallelN           7      6038430.0 ns  200275.31     64
countParallelN           8      5733226.6 ns  36285.25      64
countParallelN           9      6024557.9 ns  160247.80     64
countParallelN          10      6053407.0 ns  530306.83     64
countParallelN          11      6091350.1 ns  292242.26     64
countParallelN          12      5904037.8 ns  258931.38     64
countParallelN          13      6331139.2 ns  533701.42     64
countParallelN          14      6100342.6 ns  422783.17     64
countParallelN          15      6194005.3 ns  301002.29     64
countParallelN          16      6115450.5 ns  272703.02     64
countParallelN          17      6350013.3 ns  312250.80     64
countParallelN          18      6556236.1 ns  1034255.43    64
countParallelN          19      6401026.8 ns  235667.98     64
countParallelN          20      6335108.0 ns  276221.41     64
countParallelN          21      6403862.1 ns  436724.30     64
countParallelN          22      6254358.7 ns  108029.33     64
countParallelN          23      6486949.3 ns  403380.18     64
countParallelN          24      6484977.8 ns  391171.35     64
countParallelN          25      6402272.5 ns  177621.02     64
countParallelN          26      6587896.1 ns  453236.33     64
countParallelN          27      6671015.9 ns  549125.63     64
countParallelN          28      6539550.9 ns  237572.40     64
countParallelN          29      6756815.7 ns  382927.78     64
countParallelN          30      6776902.6 ns  440091.76     64
countParallelN          31      6764853.6 ns  409074.90     64
countParallelN          32      6823501.8 ns  403032.32     64
```

Based on the results, when the number of threads is 2, it gives the best performance, it has the lowest mean running time. This may be because that my computer has 2 cores (not as the system info, but the information on my computer). And we can also see that when the number of threads is 4, 6 and 8, they have a better performance. Maybe because they can be divided by 2 - the number of cores. But when the number of threads are bigger than 8, we can't see a significant difference.

3. Now instead of the LongCounter class, use the `java.util.concurrent.atomic.AtomicLong` class for the counts. Perform the measurements again as indicated above. Discuss the results: is the performance of AtomicLong better or worse than that of LongCounter? Should one in general use adequate built-in classes and methods when they exist?

```

# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Oracle Corporation; 17.0.3
# CPU:   null; 4 "cores"
# Date:  2022-10-14T00:02:51+0200
countSequential          9808714.3 ns  205362.00      32
countParallelN          1  9731490.7 ns  42877.05      32
countParallelN          2  6153951.4 ns  432055.12     64
countParallelN          3  6509242.1 ns  46128.94     64
countParallelN          4  5740885.5 ns  74381.42     64
countParallelN          5  6000743.4 ns  42408.12     64
countParallelN          6  5745877.0 ns  168440.64     64
countParallelN          7  5881720.6 ns  77541.73     64
countParallelN          8  5928870.8 ns  481795.17     64
countParallelN          9  5868264.4 ns  126325.55     64
countParallelN         10  5900890.9 ns  257207.72     64
countParallelN         11  5915625.4 ns  110263.80     64
countParallelN         12  5903269.2 ns  271632.53     64
countParallelN         13  6009187.0 ns  173450.31     64
countParallelN         14  6011901.6 ns  255129.99     64
countParallelN         15  6016384.9 ns  224439.55     64
countParallelN         16  6063277.4 ns  294851.70     64
countParallelN         17  6073173.1 ns  170994.97     64
countParallelN         18  6503371.9 ns  686517.34     32
countParallelN         19  6551316.0 ns  482633.60     64
countParallelN         20  6276914.9 ns  419257.23     64
countParallelN         21  6520606.5 ns  455838.68     64
countParallelN         22  6341613.7 ns  382063.24     64
countParallelN         23  6442255.3 ns  469971.44     64
countParallelN         24  6464152.7 ns  569967.67     64
countParallelN         25  6470200.0 ns  394618.30     64
countParallelN         26  6866614.5 ns  734571.52     64
countParallelN         27  6542770.6 ns  437615.06     64
countParallelN         28  6643148.4 ns  427067.34     64
countParallelN         29  6457241.5 ns  55336.02     64
countParallelN         30  6732200.1 ns  521407.05     64
countParallelN         31  6710040.1 ns  427852.87     64
countParallelN         32  7106116.3 ns  546215.39     64

```

When use AtomicLong class, the mean run time is longer.

Exercise 5.4

1. Use Mark7 (from Benchmark.java) to compare the performance of incrementing a volatile int and a normal int. Include the results in your hand-in and comment on them: Are they plausible? Any surprises?

```

# OS:   Mac OS X; 12.3.1; x86_64
# JVM:  Oracle Corporation; 17.0.3
# CPU:   null; 4 "cores"
# Date:  2022-10-14T00:48:06+0200
test normal int          1.2 ns      0.01  268435456
test volatile int        7.2 ns      0.09  67108864

```

The code is within *TestVolatile.java*, the result is shown as above. From the result we can see that the performance of incrementing a volatile int is worse than a normal int.

Volatile int always read the value from the main memory, this may result in some error, when some threads changed the value in the register but hasn't write the value into the main memory. The performance of the normal int is better may because that normal int read the value from the local register for each thread, which will cost less time than reading from the main memory.

Exercise 5.5

1. Extend LongCounter with these two methods in such a way that the counter can still be shared safely by several threads.

The implementation of this part can be seen in *LongCounter.java*, use *ReentrantLock*.

2. How many occurrences of "ipsum" is there in long-text-file.txt. Record the number in your solution.

```
Array Size: 5697
# Occurences of ipsum :1430
```

The size of the array is 5697, and occurrences of ipsum is 1430.

3. Use Mark7 to benchmark the search function. Record the result in your solution.

Test search function 12340448.0 ns 585142.47 32

```
Array Size: 5697
# Occurences of ipsum :1430
Test search function      12340448.0 ns 585142.47      32
```

4. Extend the code in TestTimeSearch with a new method